

Open Governance Standard (draft)

Governance as redesignable civilisational infrastructure

This Open Governance Standard (OGS) work begins from three first-principles assumptions:




First, Governance acts through rules.

Rules are the smallest durable units through which institutions encode permissions, constraints, obligations, and consequences.



Second, Rules are humanity's earliest automation technology.

Long before digital systems, societies used rules to pre-structure behaviour and decisions at scale.



Third, AI extends that layer rather than replacing it.

If modern automation is built on weak rule systems, failure is mechanised rather than removed.

Modern automation is only as good as the rule systems beneath it.

Why this work exists

From governance principles to operational architecture

Problem

Governance often stops at values, principles, and local tools without a coherent way to translate them into rules, tasks, controls, and accountable execution.

Implication

Institutions can state accountability, transparency, neutrality, and integrity clearly while still lacking the structure to make those properties real in live systems.

Response

OGS is a modular governance stack for rule change, shared governance language, atomic task design, and fail-safe orchestration.



**This is a
governance stack**
not a single tool



**This is
operational architecture**
not another values-only
framework

The Open Governance Standard Stack

Using the operating system analogy

The OGS forms a layered governance stack.



Project White Collar Framework
modular structure for governed AI execution.




Task Boundary Framework
atomic task model for human-AI work



Common Governance Language
shared grammar for governed reality.



Open Meta-Governance Standard
kernel for rule lifecycle and governance change



Governance debt emerges when rule metabolism cannot keep up

Why governance failure compounds under accelerating change

Governance fails not only when rules are wrong, but when rule systems cannot adapt fast enough.


Governance debt builds when institutions cannot update, interpret, implement, and retire rules at the speed change demands.

The issue is metabolic capacity not only rule quality.

Why this matters

When governance cannot keep up, systems remain formally governed but practically misaligned.

When rule intake exceeds rule retirement capacity, governance debt compounds.



Open Meta-Governance Standard (OMGS)

How rule changes over time


The **OMGS** focuses on the operational reality of governance: the set of rules of how rules should be administrated.

It is built from two first-principles lenses:



Lifecycle of rule change

how rules are proposed, implemented, communicated, enforced, revised, and retired



Ecosystem in which rules operate

who is governed, who decides, what resources are involved, and how rules shape reality through institutions and rule users.

OMGS makes rule systems more explicit, inspectable, and adaptable under accelerating change.

Common Governance Language (CGL)

A shared language for governed reality

CGL lets policy, operations, compliance, and technical teams describe the same governed system in compatible terms.


Why it matters

CGL's shared grammar reduces translation loss across law, policy, procedure, workflow design, and implementation.

Interoperability

It works **vertically** across law, policy, and procedure, and **horizontally** across sectors such as social services, health, and foreign affairs.

This is the language layer of the OGS stack.



Task Boundary Framework (TBF)

The atomic task is the real unit of AI governance

TBF makes governable work visible

showing where reasoning sits, what kind of judgement is involved, what information is being handled, and where decision must remain human.


It distinguishes both reasoning and information structure

including what / how / why reasoning, and more structured versus more unstructured information environments.

Why this matters

Institutions can define automation boundaries, human review points, and future work redesign more explicitly.

If AI absorbs more “what” and “how” work, organisations will need to redesign roles towards judgement, escalation, and “why” functions ahead of time.



Every information transformation = a decision

Project White Collar Framework (PWC)

Fail-safe orchestration for AI-heavy workflows

Embeds governance into live execution

through role separation, escalation paths, fail-safe controls, and reconstructable accountability.

Remains interoperable with human governance

by embedding institutional safeguards familiar from well-designed organisations.


Neutrality by design

Default approval and default rejection are both built-in bias positions. A more neutral system depends on separation of powers, diversity of perspectives, and collective decision making.

Diversity of perspective by reasoning role

Real perspective diversity comes from separating distinct reasoning function tied to the judgement outcome. Three similar AI operators do not create real diversity if they think through the same frame.

Execution without institutional design is not governance maturity.



How the layers work together

From rule formation to live execution

OGS is designed to operationalise governed reality from first principles.



Rules, descriptions, tasks, and execution are treated as parts of one architecture.

Four grounding insights

The design commitments beneath OGS



Governance fails less from lack of values than from **lack of operational structure**.



Rule quality is not enough; **rule metabolism** determines whether governance can **keep up with change**.



Governance should **optimise for long-term stewardship**, not short-term execution speed.



AI governance must remain **interoperable with human governance**.

AI operates in unstructured reasoning environments where conventional digital rules do not apply natively.

OGS is built to operationalise accountability, transparency, integrity, and neutrality.

Appendix A: Framework materials

Deeper materials for each layer of the stack

This deck is a compact overview of the stack. Further information is available in the GitHub links below.



OMGS GitHub (CC-BY-4.0)

- [Pitch deck](#)
- [Framework materials](#)



CGL GitHub (CC-BY-4.0)

- [Pitch deck](#)
- [Framework materials](#)



TBF GitHub (CC-BY-4.0)

- [Pitch deck](#)
- [Framework materials](#)



PWC GitHub (CC-BY-4.0)


- [Pitch deck](#)
- [Framework materials](#)

Published under CC BY 4.0 for reuse, adaptation, and community contribution.

Appendix B: Rule Lifecycle

Open Meta-Governance Standard


A full rule lifecycle runs from changes in the rule environment through decision, implementation, communication, and enforcement.



Appendix C: Rule Ecosystem


Open Meta-Governance Standard

OMGS also considers the ecosystem in which rules operate and how that shapes interpretation, application, and revision.



Appendix D: Enforcement mechanism as a governance design dimension

Common Governance Language



Enforcement should be designed to achieve the intended rule outcome when voluntary compliance mechanisms fail.

Appendix E: Separation of Powers and Diversity of Perspective by Design

Project White Collar

The **top accountability bot roles in PWC** are held separately so that compliance, escalation, and control are not collapsed into one execution path.



Governance Lead

Instruction Auditor

- Checks bots follow system prompts and **role boundaries**
- Can **trigger independent review** or veto workflow

Defending
system integrity



Escalator

Risk Sentinel

- Detects **ethical and moral concerns** beyond routine compliance
- Escalates issues **outside normal execution** flow

Protecting owner and
user interests



Controller

Blindfolded Judge

- Turns bots **on/off, mutes**, or suspends execution
- **Grants or denies access** to search, drive, APIs, and other system resources

Mechanical
authority only

Distributed oversight without overreach is a design choice, not an afterthought.

Appendix F: Operationalising fail-safe AI orchestration

Project White Collar

The point is not maximum efficiency, but governed performance under conditions of consequence, failure, and sustained operation.



One-shot orchestration

Supercar mode



Fail-safe orchestration

Standard ship mode

Design goal	Speeding as fast as possible.	Carry value safely, for as long as needed
Speed brag	0 – 100: 2.2 seconds!	0 – 100: I don't care.
Sustained operation	Max rev: minutes before something melts.	Max rev: weeks; limited by fuel and maintenance.
Failure assumption	We assume it usually works. Failure = accident.	We assume components fail. Surviving failure = a design case.
When something goes wrong	"Oh no!" Pull over, call a tow truck.	"Use the SOPs." Redundancy, crew, reroute + repair underway.
Where you use it	Weekend thrills, racetracks, low-duty use.	Cargo, passengers, critical routes.

If your agents move money, change records, or trigger real-world actions, you are running a ship, not a supercar.